

Minimum Cut, Bipartite Matching and Minimum Cost Maximum Flow with BGL

Xun Zou¹

ETH Zürich

November 17, 2021

¹based on material from Daniel Graf, Andreas Bärtschi

Recap: Basic Network Flows – What did we see last time?

Maximum Flow

- ▶ Edge-disjoint paths
- ▶ Circulation Problem
- ▶ Flow applications

Common Tricks

- ▶ Multiple sources/targets
- ▶ Undirected edges
- ▶ Vertex capacity
- ▶ Minimum flow constraint

Today: Advanced Network Flows – What else are flows useful for?

Cuts in directed graphs

- ▶ How to disconnect one vertex from another by deleting edges?

Bipartite matchings

- ▶ How to compute a maximum matching in a bipartite graph with flow?


Flows with Costs

- ▶ What if sending flow comes with a price?

Minimum Cut: Shopping Trip



Minimum Cut: Shopping Trip



Start from HB:

- ▶ Visit as many shops as possible.
- ▶ Return to HB after each shop.

Condition: Use each road in at most one trip.

Result: the number of **edge-disjoint paths** \Rightarrow 4 shops.

- ▶ The **bottleneck** between s and t .

Unrealistic condition!

(There are interesting streets in Zürich.)

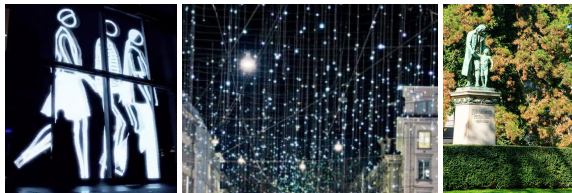
Minimum Cut: Shopping Trip



Start from HB:

- ▶ Visit as many shops as possible.
- ▶ Return to HB after each shop.

Condition: **Beautiful roads may be used more than once.**



We may use Bahnhofstrasse up to three times.

Result: the **weighted bottleneck** \Rightarrow 6 shops.

- ▶ minimum cut between s and t .

Minimum Cut: Cuts and Flows

$G = (V, E, s, t)$ a flow network.


► $S \subset V$ s.t. $s \in S, t \in V \setminus S (=: T)$, e.g. $S = \{s, v_1, v_3\}$

The value of the (S, T) -cut is

$\text{cap}(S, T) :=$ outgoing capacity

$$= \sum_{\substack{e=(u,v) \\ u \in S, v \in T}} \text{cap}(e)$$

$$= 3 + 2 = 5.$$




The value of a flow f from S to T is

$f(S, T) :=$ outgoing flow – incoming flow

$$= \sum_{\substack{e=(u,v) \\ u \in S, v \in T}} \text{flow}(e) - \sum_{\substack{e=(v,u) \\ v \in T, u \in S}} \text{flow}(e)$$

$$= 1 + 2 - 1 = 2.$$



Minimum Cut: Maxflow-Mincut-Theorem

Theorem (Maxflow-Mincut-Theorem)

Let f be an s - t -flow in a graph G . Then f is a maximum flow if and only if

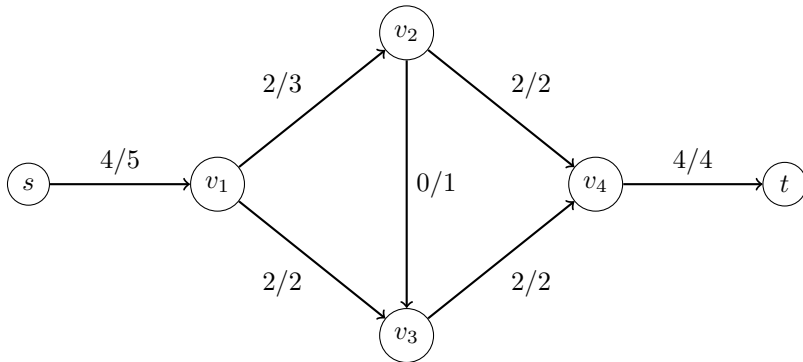
$$|f| = \min_{S: s \in S, t \notin S} \text{cap}(S, V \setminus S).$$

This allows us to easily find a minimum s - t -cut:

- ▶ Construct the residual graph $G_f := (V, E_f)$. For each edge $(u, v) \in G$ we have:
 - An edge $(u, v) \in G_f$ with capacity $\text{cap}(e) - f(e)$, if $\text{cap}(e) - f(e) > 0$.
 - An edge $(v, u) \in G_f$ with capacity $f(e)$, if $f(e) > 0$.
- ▶ Since f is a maximum flow, there is no s - t path in the residual graph G_f .
- ▶ Take S to be all vertices in G_f reachable from s .
 $\Rightarrow (S, V \setminus S)$ is a minimum s - t -cut.

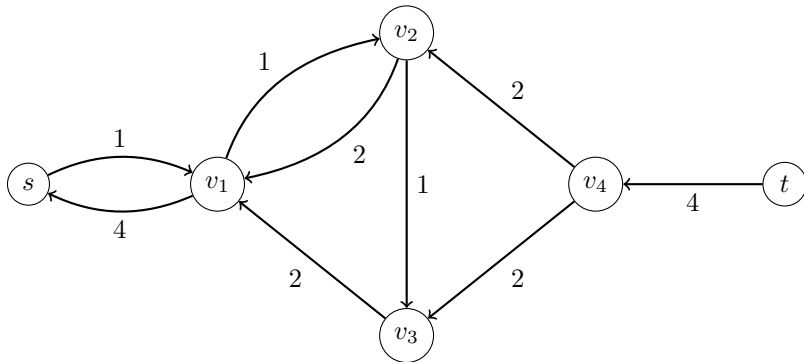
Minimum Cut: Example

Graph G and a maximum flow f .



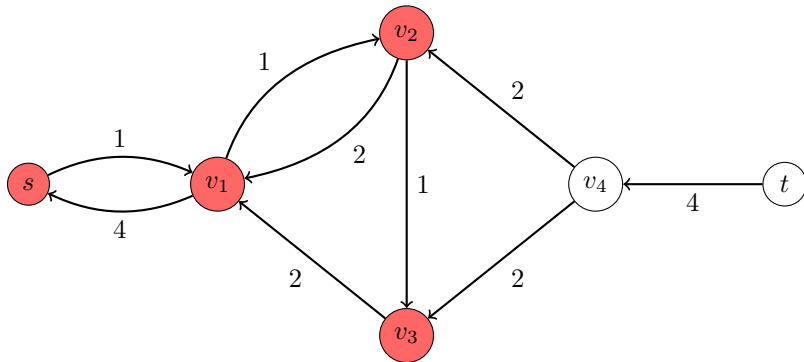
Minimum Cut: Example

Residual graph G_f .



Minimum Cut: Example

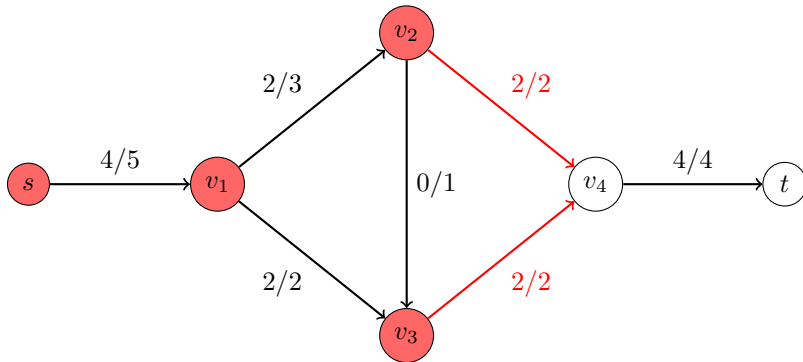
Residual graph G_f .



$$S = \{s, v_1, v_2, v_3\}$$

Minimum Cut: Example

Residual graph G_f .



$$S = \{s, v_1, v_2, v_3\} \Rightarrow \text{cap}(S, V \setminus S) = |f| = 4.$$

Minimum Cut: Proof of the Maxflow-Mincut-Theorem

Theorem (Maxflow-Mincut-Theorem)

Let f be an s - t -flow in a graph G . Then the following are equivalent:

1. $|f| = \min_{s \in S, t \notin S} \text{cap}(S, V \setminus S)$.
2. f is a maxflow.
3. There is no s - t path in the residual graph G_f .

Proof.

- 1) \implies 2) f cannot be bigger than $\min_S \text{cap}(S, V \setminus S)$.
- 2) \implies 3) Indirectly: If there was s - t path in G_f , f could be extended.
- 3) \implies 1) Take S to be all vertices in G_f reachable from s .
Then all edges from S to $V \setminus S$ must be fully saturated by the flow, and the incoming flow to S must be 0.
But then $|f| = f(S, V \setminus S) = \text{cap}(S, V \setminus S)$.



Minimum Cut: Code

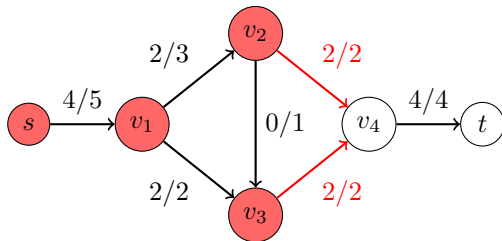
Example code: BFS on the residual graph G_f . See `bgl_residual_bfs.cpp` on moodle.

```
79 // BFS to find vertex set S
80 std::vector<int> vis(N, false); // visited flags
81 std::queue<int> Q; // BFS queue (from std:: not boost::)
82 vis[src] = true; // Mark the source as visited
83 Q.push(src);
84 while (!Q.empty()) {
85     const int u = Q.front();
86     Q.pop();
87     OutEdgeIt ebegin, eend;
88     for (boost::tie(ebegin, eend) = boost::out_edges(u, G); ebegin != eend; ++ebegin) {
89         const int v = boost::target(*ebegin, G);
90         // Only follow edges with spare capacity
91         if (rescapacitymap[*ebegin] == 0 || vis[v]) continue;
92         vis[v] = true;
93         Q.push(v);
94     }
95 }
```

Minimum cut: Algorithm

Summary of what you need to do to find a minimum cut:

1. Compute maximum flow f and the residual graph G_f .
2. Compute the set S of vertices that are reachable from the source s in G_f .
 - BFS on edges with residual capacity > 0 .
3. Output (depending on the task):
 - All vertices in S .
 - All edges going from S to $V \setminus S$.

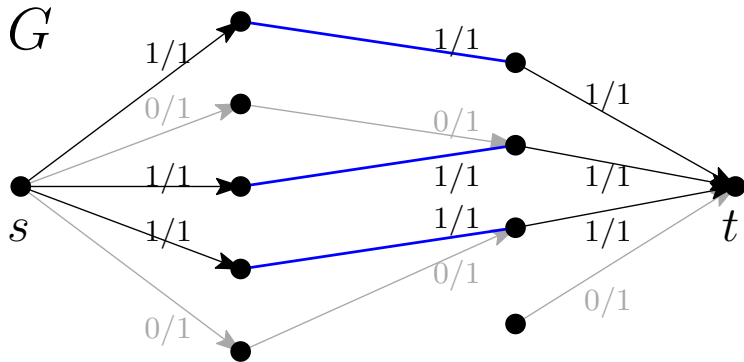


Note: minimum cuts are not necessarily unique, but the earliest one is.

Maximum Matchings: Bipartite Graphs

Maximum Matching: pick as many non-adjacent edges as possible

Flow formulation through vertex capacities / edges-disjoint paths:



Vertex Cover and Independent Set: General Graphs

- ▶ **Maximum independent set (MaxIS)**

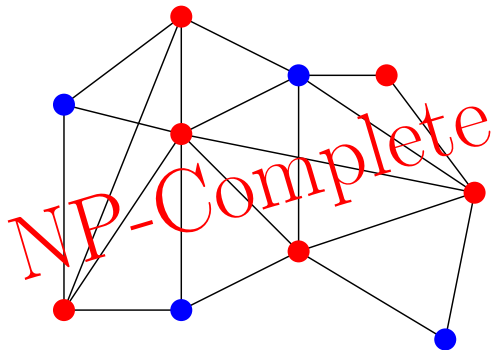
Largest $I \subseteq V$, such that
 $\nexists u, v \in I : (u, v) \in E$.

- ▶ **Minimum vertex cover (MinVC)**

Smallest $C \subseteq V$, such that
 $\forall (u, v) \in E : u \in C \vee v \in C$.

- ▶ These problems are complementary!

$$\text{MaxIS} = V \setminus \text{MinVC}$$



Vertex Cover and Independent Set: Bipartite Graphs

Theorem (König: MinVC and MaxIS are simpler on bipartite graphs!)

In a bipartite graph, the number of edges in a maximum matching is equal to the number of vertices in a minimum vertex cover.

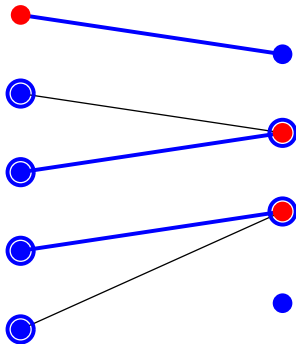
Proof: See [Wikipedia](#) for a nice and short proof.

Algorithm:

1. Maximum matching M , $V = L \cup R$. Find all unmatched vertices in L , label them as visited.
2. Starting at visited vertices search (BFS) left to right along edges from $E \setminus M$ and right to left along edges from M . Label each found vertex as visited.
3. MinVC – all unvisited in L and all visited in R .
MaxIS – all visited in L and all unvisited in R .

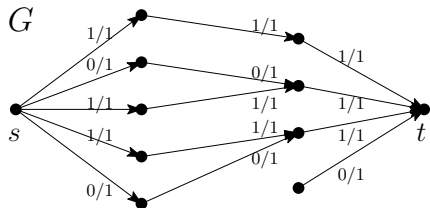
Careful! Step 2 can take several rounds.

Easy Implementation?




Finding a MinVC or MaxIS in bipartite graphs: step by step

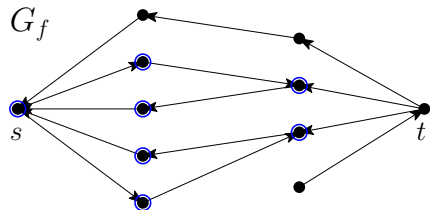
1) Formulate and compute the flow:




2) Compute the residual graph G_f :



3) Mark reachable vertices from s with BFS:



4) Read the MinVC or MaxIS from the marks:



Summary: MaxFlowMinCut and Bipartite Matching

What you should remember:

Edge Cut

- ▶ Theorem: maximum amount of any s - t -flow = minimum capacity of any s - t -cut
- ▶ Finding the cut: BFS/DFS on residual graph starting from s .

Vertex Cover

- ▶ Minimum vertex cover and maximum independent set are hard problems.
- ▶ Bipartite graphs allow fast MinVC and MaxIS (both on top of maximum matching).
- ▶ Finding the minimum vertex cover: BFS/DFS on residual graph from s .

Minimum Cost for a Bipartite Matching


What if the edges in a matching also have costs associated?

- ▶ cardinality of the matching is no longer the only objective
- ▶ second priority: minimize the total cost


Searching for the cheapest among all maximum matchings?

Need for new tools.

two maximum matchings of different cost:



$$1 + 5 + 3 = 9$$



$$2 + 1 + 2 = 5$$


More General Model: Minimum Cost Maximum Flow

Input: a flow network consisting of

- ▶ a directed graph $G = (V, E)$
- ▶ a source and a sink $s, t \in V$
- ▶ edge capacity $cap : E \rightarrow \mathbb{N}$
- ▶ **edge cost** $cost : E \rightarrow \mathbb{Z}$.

Output: a flow f with minimal
 $cost(f) = \sum_{e \in E} f(e) \cdot cost(e)$
among all flows with maximal $|f|$.

Note: it can model much more than just
minimum cost bipartite matching.



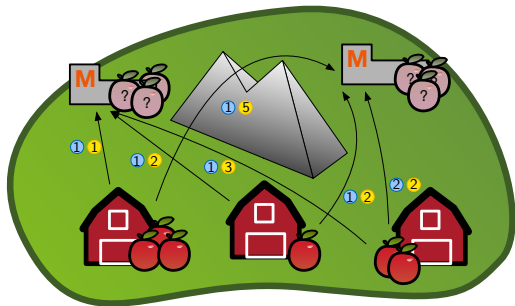
Legend: capacity ● ● cost

Example: Fruit Delivery

A supermarket wants to schedule fruit deliveries from their farmers to their shops.


They know all important parameters:

- ▶ production per farm [in kg]
- ▶ demand per shop [in kg]
- ▶ transportation capacity [in kg] and transportation cost [in CHF pro kg] for every farm-shop pair




Note: this is not just a bipartite matching, even though the graph is bipartite. One farm might deliver to multiple shops (and vice versa).

Example: Fruit Delivery




$$\text{Flow: } 2 + 1 + 2 = 5$$
$$\text{Cost: } 1 \cdot 1 + 1 \cdot 5 + 1 \cdot 2 + 2 \cdot 2 = 12$$

Example: Fruit Delivery



Extended Example: Canned Fruit Delivery

Extension: canned fruit requires transportation to and from a canning factory.




Min Cost Max Flow with BGL

There are two algorithms available in BGL (available in BGL v1.55+):

- ▶ `cycle_canceling()`
 - ▶ slow, but can handle negative costs
 - ▶ needs a maximum flow to start with (call e.g. `push_relabel_max_flow` before)
 - ▶ runtime $\mathcal{O}(C \cdot (nm))$ where C is the cost of the initial flow
 - ▶ [\[BGL documentation\]](#), [\[BGL example\]](#).
- ▶ `successive_shortest_path_nonnegative_weights()`
 - ▶ faster, but works only for non-negative costs
 - ▶ sum up all residual capacities at the source to get the flow value
 - ▶ runtime $\mathcal{O}(|f| \cdot (m + n \log n))$
 - ▶ [\[BGL documentation\]](#), [\[BGL example\]](#).

Rough guide for $m \approx n$, $|C|, |f| \ll n$



Min Cost Max Flow with BGL

Weights and capacities, just one more nesting level in the typedefs:

```
16 // Graph Type with nested interior edge properties for Cost Flow Algorithms
17 typedef boost::adjacency_list_traits<boost::vecS, boost::vecS, boost::directedS> traits;
18 typedef boost::adjacency_list<boost::vecS, boost::vecS, boost::directedS, boost::no_property,
19     boost::property<boost::edge_capacity_t, long,
20     boost::property<boost::edge_residual_capacity_t, long,
21     boost::property<boost::edge_reverse_t, traits::edge_descriptor,
22     boost::property<boost::edge_weight_t, long> > > > graph; // new! weightmap corresponds to
23
24 typedef boost::graph_traits<graph>::edge_descriptor      edge_desc;
25 typedef boost::graph_traits<graph>::out_edge_iterator    out_edge_it; // Iterator
```

Code file: bgl_mincostmaxflow.cpp

Min Cost Max Flow with BGL

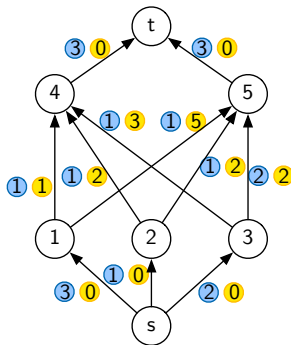
Extending the edge adder:

```
27 // Custom edge adder class
28 class edge_adder {
29     graph &G;
30
31 public:
32     explicit edge_adder(graph &G) : G(G) {}
33     void add_edge(int from, int to, long capacity, long cost) {
34         auto c_map = boost::get(boost::edge_capacity, G);
35         auto r_map = boost::get(boost::edge_reverse, G);
36         auto w_map = boost::get(boost::edge_weight, G); // new!
37         const edge_desc e = boost::add_edge(from, to, G).first;
38         const edge_desc rev_e = boost::add_edge(to, from, G).first;
39         c_map[e] = capacity;
40         c_map[rev_e] = 0; // reverse edge has no capacity!
41         r_map[e] = rev_e;
42         r_map[rev_e] = e;
43         w_map[e] = cost;    // new assign cost
44         w_map[rev_e] = -cost; // new negative cost
45     }
46 };
```

Min Cost Max Flow with BGL

Building the graph:

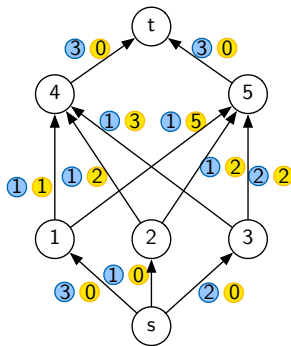
```
50     const int N=7;
51     const int v_source = 0;
52     const int v_farm1 = 1;
53     const int v_farm2 = 2;
54     const int v_farm3 = 3;
55     const int v_shop1 = 4;
56     const int v_shop2 = 5;
57     const int v_target = 6;
58
59     // Create graph, edge adder class and property maps
60     graph G(N);
61     edge_adder adder(G);
62     auto c_map = boost::get(boost::edge_capacity, G);
63     auto r_map = boost::get(boost::edge_reverse, G);
64     auto rc_map = boost::get(boost::edge_residual_capacity, G);
```



Min Cost Max Flow with BGL

Add the edges:

```
67 adder.add_edge(v_source, v_farm1, 3, 0);
68 adder.add_edge(v_source, v_farm2, 1, 0);
69 adder.add_edge(v_source, v_farm3, 2, 0);
70
71 adder.add_edge(v_farm1, v_shop1, 1, 1);
72 adder.add_edge(v_farm1, v_shop2, 1, 5);
73 adder.add_edge(v_farm2, v_shop1, 1, 2);
74 adder.add_edge(v_farm2, v_shop2, 1, 2);
75 adder.add_edge(v_farm3, v_shop1, 1, 3);
76 adder.add_edge(v_farm3, v_shop2, 2, 2);
77
78 adder.add_edge(v_shop1, v_target, 3, 0);
79 adder.add_edge(v_shop2, v_target, 3, 0);
```



Min Cost Max Flow with BGL

Running the algorithm:

```
84 // Option 1: Min Cost Max Flow with cycle_canceling
85 int flow1 = boost::push_relabel_max_flow(G, v_source, v_target);
86 boost::cycle_canceling(G);
87 int cost1 = boost::find_flow_cost(G);
88 std::cout << "-----" << "\n";
89 std::cout << "Minimum Cost Maximum Flow with cycle_canceling()" << "\n";
90 std::cout << "flow " << flow1 << "\n"; // 5
91 std::cout << "cost " << cost1 << "\n"; // 12
```


Min Cost Max Flow with BGL

Running the algorithm:

```
92 // Option 2: Min Cost Max Flow with successive_shortest_path_nonnegative_weights
93 boost::successive_shortest_path_nonnegative_weights(G, v_source, v_target);
94 int cost2 = boost::find_flow_cost(G);
95 std::cout << "-----" << std::endl;
96 std::cout << "Minimum Cost Maximum Flow with successive_shortest_path_nonnegative_weights()" << "\n";
97 std::cout << "cost " << cost2 << "\n"; // 12
98 // Iterate over all edges leaving the source to sum up the flow values.
99 int s_flow = 0;
100 out_edge_it e, eend;
101 for(boost::tie(e, eend) = boost::out_edges(boost::vertex(v_source,G), G); e != eend; ++e) {
102     s_flow += c_map[*e] - rc_map[*e];
103 }
104 std::cout << "s-out flow " << s_flow << "\n"; // 5
105 // Or equivalently, you can do the summation at the sink, but with reversed edge.
106 int t_flow = 0;
107 for(boost::tie(e, eend) = boost::out_edges(boost::vertex(v_target,G), G); e != eend; ++e) {
108     t_flow += rc_map[*e] - c_map[*e];
109 }
110 std::cout << "t-in flow " << t_flow << "\n"; // 5
```

Summary: Min Cost Max Flow with BGL

What you should remember from this part:

Minimum Cost Maximum Flow

- ▶ is a powerful and versatile modeling tool.
- ▶ is a tiebreaker among several maximum flows (but might still not be unique).
- ▶ = *maximum* cost maximum flow with negated costs.
- ▶ can often be reformulated without negative costs which allows us to use a faster algorithm in BGL (key step in many problems).
- ▶ can easily be implemented when starting with our template on Moodle.

If you are interested in the theory behind these algorithms: (not needed for this course)

- ▶ Stanford CS 261 by Prof. Tim Roughgarden, Spring 2016, [full lectures on Youtube](#)

A little bit of history...

What we do in Algotlab today was cutting edge research in 1955 ;-)

→ *“On the history of the transportation and maximum flow problems”*
by Alexander Schrijver, *Mathematical Programming*, 2002

A little bit of conclusion...

No more new theory and tools past this point.

But be prepared to combine all your skills:

- ▶ On top of a flow problem, do binary search for the answer
- ▶ LP formulation vs. flow formulation?
- ▶ Some graph problems can be solved greedily (e.g. MST), others not (e.g. flow)
- ▶ Find a Min Cost Max Flow formulation where greedy fails for non-unit weights
- ▶ Do BFS on Delaunay triangulation or do Union-Find on Euclidean MST
- ▶ ...

Starting next week:

- ▶ How to balance reading, solving, coding, debugging under time constraints
- ▶ No more problems labeled by topic – figure it out yourself
- ▶ In-class exercises on Wednesdays